

Pavan Kartheek Konam

AI Engineer | Chicago, IL | pavankonam@gmail.com | +1 (517) 940-2499 | [linkedin.com/in/pavankonam](https://www.linkedin.com/in/pavankonam) | github.com/pavankonam | [portfolio](#) | ruvaah.com

PROFESSIONAL SUMMARY

AI engineer with an M.S. in Artificial Intelligence from Northwestern. Focus areas: LLMs, RAG, and multi-agent systems. Currently Founding Engineer at Ruvaah – live at ruvaah.com – flagging hiring fraud before the interview, and open to full-time roles. Before that, several years building ML systems across healthcare, e-commerce, and analytics. Comfortable owning a feature end to end: UI, vector store, fine-tuned models. Builds fast with Claude Code and Cursor.

EXPERIENCE

Founding Engineer | *Ruvaah* (ruvaah.com) Apr 2026 – Present

Human Trust Intelligence platform that flags hiring fraud before the interview, scoring candidates in real time across nine signal tracks.

- Took the product from zero to its first paying customers in weeks, architecting the full stack solo: a Next.js frontend, a Fastify/TypeScript API, and a Python FastAPI ML service over PostgreSQL, Neo4j, and Redis.
- Built the nine-track scoring engine behind the real-time trust score – document forensics, GitHub-activity analysis, and graph-based fraud-ring detection over Neo4j – to catch clusters of linked or recycled identities.
- Designed multi-tenant architecture with role-based score weighting so each customer's data and scoring stay fully separate; iterated quickly with Claude Code and Cursor.

AI Software Engineer (Summer Internship) | *HauteCarat* Jul 2025 – Sep 2025

Luxury jewelry e-commerce, established scale-up.

- Built Intellicore, a RAG assistant over a Pinecone-backed knowledge layer of product catalogs, sales, and customer feedback, so leadership and the marketing, design, and operations teams could query live business data in plain English (OpenAI, LangChain, FastAPI) – cut analysis and reporting from weeks to days.
- Fine-tuned LLMs and embedding models on domain data with Hugging Face Transformers, built the data pipelines feeding them (Pandas, Airflow), and shipped a computer-vision pipeline (PyTorch, OpenCV, YOLOv8) that turned 3D renders into on-model product imagery, all wired into CI/CD with Docker, GitHub Actions, and MLflow.

Graduate Teaching Assistant & Lecturer | *Northwestern University* Jan 2025 – Dec 2025

- TA'd two Northwestern ML and AI courses through 2025 and independently designed and taught a summer linear-algebra and MATLAB course, translating complex concepts into hands-on labs for ~30 engineering students in the EXCEL program.

Cognizant Feb 2022 – Aug 2024

Programming Analyst – AI/ML Jul 2022 – Aug 2024

- Increased clinical efficiency over 25% by architecting an AI model for CVS Health that streamlined patient-data processing and medical alerts.
- Hit 92%+ production accuracy on a clinical speech-to-text system by building the core PyTorch and TensorFlow models, with SQL ETL pipelines, model versioning, and cloud inference on AWS and Azure behind it.
- Lifted model performance about 20% through feature engineering and statistical analysis on large, messy clinical datasets, and built the ML pipelines and quality checks that kept them reliable in a regulated healthcare setting.

Computer Vision & ML Intern Feb 2022 – Jun 2022

- Trained computer-vision models in PyTorch (~94% accuracy on medical object detection) and deployed them into diagnostics workflows on AWS, with Azure and GCP, working with real clinical data from training through deployment.

EDUCATION

Northwestern University Dec 2025 *M.S. in Artificial Intelligence, GPA: 3.85/4.00*

Certifications: Google Cloud – Data Engineering, Cloud Architecture

PROJECTS

ICU Insights – Predictive ML on Clinical Data Sep 2024 – Dec 2024

- Predicted ICU admissions from 70,000+ MIMIC-IV patient records with Random Forest and neural-net models (AUC-ROC 0.973, F1 0.767), feeding them through BigQuery ETL pipelines with chunked processing and engineered features.

Multi-Agentic Planner – Agent Orchestration on Azure Jan 2025 – Mar 2025

- Stood up a four-agent pipeline (planner, retriever, executor, reviewer) on Azure Durable Functions to automate end-to-end task planning and execution, coordinating agents through function-calling and RAG with fallback handling and observability built in.

Medicurious – Domain Healthcare LLM Jan 2025 – Mar 2025

- Grounded a healthcare LLM in curated clinical knowledge (RAG, vector embeddings, OpenAI/Anthropic APIs) to cut hallucinations, and wrote eval protocols scoring answers on accuracy and clinical relevance.

TECHNICAL SKILLS

- **AI & ML:** LLMs, RAG, Multi-Agent Systems, Fine-tuning, Embeddings, Prompt Engineering, Evaluation, LangChain, LangGraph, NLP, Computer Vision, GANs, Deep Learning, Hugging Face, PyTorch, TensorFlow, OpenAI & Anthropic APIs
- **Languages:** Python, TypeScript, SQL, C/C++, R, MATLAB
- **Backend & Data:** FastAPI, Fastify, Next.js, PostgreSQL, Neo4j, Redis, Vector Databases (Pinecone, Weaviate, ChromaDB), Pandas, scikit-learn
- **Cloud & MLOps:** AWS (SageMaker, EMR), Azure (Azure ML, Azure OpenAI), GCP, BigQuery, Snowflake, Docker, GitHub Actions, CI/CD, MLflow, Airflow